

OBL2-OS

August 12, 2020

This is a mandatory assignment. Use resources from the course to answer the following questions. **Take care to follow the numbering structure of the assignment in your submission.** Some questions may require a little bit of web searching. Some questions require you to have access to a Linux machine, for example running natively or virtually on your own PC, or by connecting to `gremlin.stud.iie.ntnu.no` over SSH (Secure Shell). Working in groups is **permitted**, but submissions must be **individual**.

1 Processes and threads

1. Explain the difference between a process and a thread.
2. Describe a scenario where it is desirable to:
 - Write a program that uses threads for parallel processing
 - Write a program that uses processes for parallel processing
3. Explain why each thread requires a thread control block (TCB).
4. What is the difference between cooperative (voluntary) threading and pre-emptive (involuntary) threading? Briefly describe the steps necessary for a context switch for each case.

2 C program with POSIX threads

*nix operating systems use POSIX threads, which are provided by the `pthread` library. Consider the following adapted code from the textbook (the code has been modified slightly to use `pthread`, while the book assumes its own thread implementation).

```
#include <stdio.h>
#include <pthread.h>
#define NTHREADS 10
pthread_t threads[NTHREADS];
void *go (void *n) {
    printf("Hello from thread %ld\n", (long)n);
    pthread_exit(100 + n);
    // REACHED?
}

int main() {
    long i;
    for (i = 0; i < NTHREADS; i++) pthread_create(&threads[i], NULL, go, (void*)i);
    for (i = 0; i < NTHREADS; i++) {
        long exitValue;
        pthread_join(threads[i], (void*)&exitValue);
```

```

        printf("Thread %ld returned with %ld\n", i, exitValue);
    }
    printf("Main thread done.\n");
    return 0;
}

```

We can compile the code and tell the compiler to link the `pthread` library:

```
$ gcc -o threadHello threadHello.c -lpthread
```

At the command prompt, run the program using `./threadHello`. The program gives output similar to the following:

```

Hello from thread 0
Hello from thread 3
Hello from thread 5
Hello from thread 1
Hello from thread 4
Thread 0 returned with 100
Thread 1 returned with 101
Hello from thread 9
Hello from thread 8
Hello from thread 2
Hello from thread 7
Hello from thread 6
Thread 2 returned with 102
Thread 3 returned with 103
Thread 4 returned with 104
Thread 5 returned with 105
Thread 6 returned with 106
Thread 7 returned with 107
Thread 8 returned with 108
Thread 9 returned with 109
Main thread done.

```

Study the code and the output. Run the code several times. Answer the following questions.

1. Which part of the code (e.g., the task) is executed when a thread runs? Identify the function and describe briefly what it does.
2. Why does the order of the “Hello from thread X” messages change each time you run the program?
3. What is the *minimum* and *maximum* number of threads that could exist when thread 8 prints “Hello”?
4. Explain the use of `pthread_join` function call.
5. What would happen if the function `go` is changed to behave like this:

```

void *go (void *n) {
    printf("Hello from thread %ld\n", (long)n);
    if(n == 5)
        sleep(2); // Pause thread 5 execution for 2 seconds
    pthread_exit(100 + n);
    // REACHED?
}

```

6. When `pthread_join` returns for thread X, in what state is thread X?